

COMP 4511

Process

Goal

- ▶ Overview of Process
- ▶ Hands-on practices via POSIX
- ▶ Lab Exercise: Create a simple shell with child process creation




Overview of Process

What is a process?

- ▶ Definition: an executable instance of a program
 - ▶ A process is the context maintained for an executing problem
 - ▶ Process is different from a program
- ▶ Process provides each program with **two key abstractions**:
 - ▶ Logical control flow
 - ▶ Each program seems to have exclusive use of the processor
 - ▶ Private virtual address
 - ▶ Each program seems to have exclusive use of main memory
- ▶ **How are these illusions maintained?**
 - ▶ Processes executions interleaved (multitasking) or run on separate cores (parallel)
 - ▶ Private address spaces managed by virtual memory system

Context switching

- ▶ Processes are managed by the kernel
- ▶ Control passes from one process to another via a context switch



What makes up a process?

- ▶ Program code
- ▶ Machine registers
- ▶ Global data
- ▶ Stack
- ▶ Open files
- ▶ An environment

Process context

- ▶ Process ID (`pid`)
- ▶ Parent process ID (`ppid`)
- ▶ Current directory
- ▶ File descriptor table
- ▶ Environment
- ▶ Pointer to program code
- ▶ Pointer to data (memory for global variables)
- ▶ Pointer to stack (memory for local variables)
- ▶ Pointer to heap (dynamically allocated memory)
- ▶ Execution priority
- ▶ Signal information

```
pid_t myid = getpid();  
pid_t myparentid = getppid();
```

Linux processes

- ▶ Virtual address space
 - ▶ The virtual address space is the memory that contains the code to execute as well as the process stack and data
- ▶ Process descriptor (`struct task_struct`): data structure in the kernel to keep track of that process
 - ▶ Virtual address space map
 - ▶ Current status of the process
 - ▶ Execution priority of the process
 - ▶ Resource usage of the process
 - ▶ Current signal mask
 - ▶ Owner of the process

Hands-on practices via POSIX


What is POSIX?

- ▶ An acronym for Portable Operating System Interface
- ▶ POSIX defines the application programming interface or software compatibility with Unix-like operating systems and other operating systems
- ▶ On Unix-like systems, unistd.h provides access to the POSIX for C/C++ programming
 - ▶ It provides process-related function calls such as `fork()`, `pipe()`,...
 - ▶ They are wrapper functions for the related system calls in Linux kernel such as `sys_fork()`, `sys_pipe()`, ...
 - ▶ Reference: <http://pubs.opengroup.org/onlinepubs/9699919799/basedefs/unistd.h.html>

Creating a process: fork

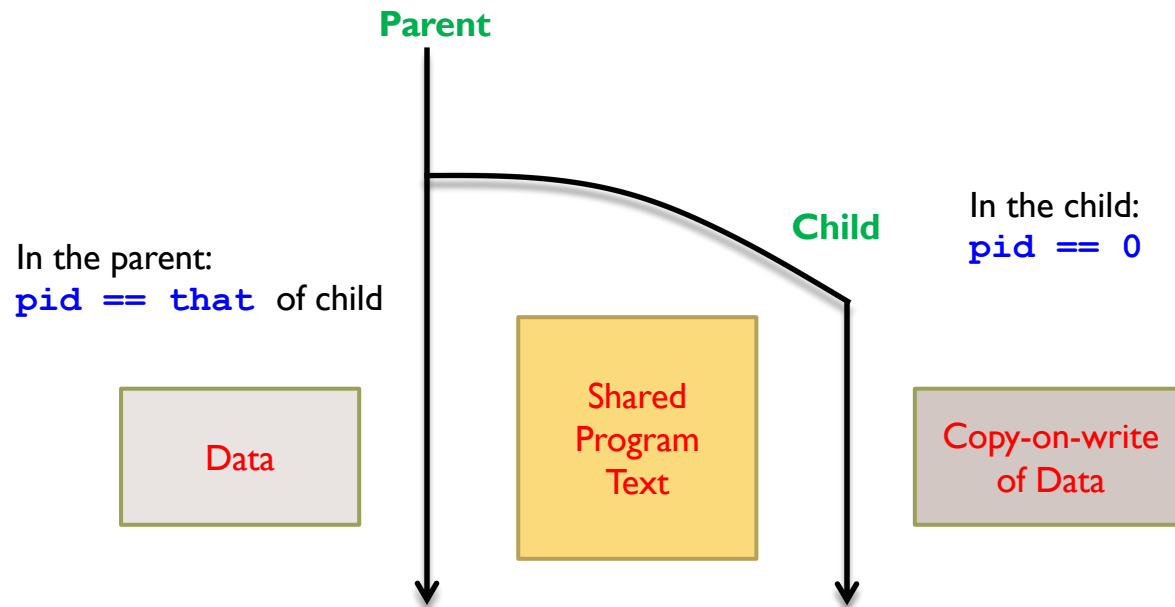
```
#include <unistd.h>
#include <sys/types.h>
pid_t fork(void);
```

- ▶ Create a child process
 - ▶ The child is an (almost) exact copy of the parent
 - ▶ The new process and the old process both continue in parallel from the statement that follows the `fork()`
- ▶ Returns
 - ▶ To child:
 - ▶ 0 on success
 - ▶ To parent:
 - ▶ Process ID of the child process
 - ▶ -1 on error, sets `errno`



The return PID is important for both parent and child processes

Creating a process: fork()



`fork()` is called **once**, but returns **twice**!

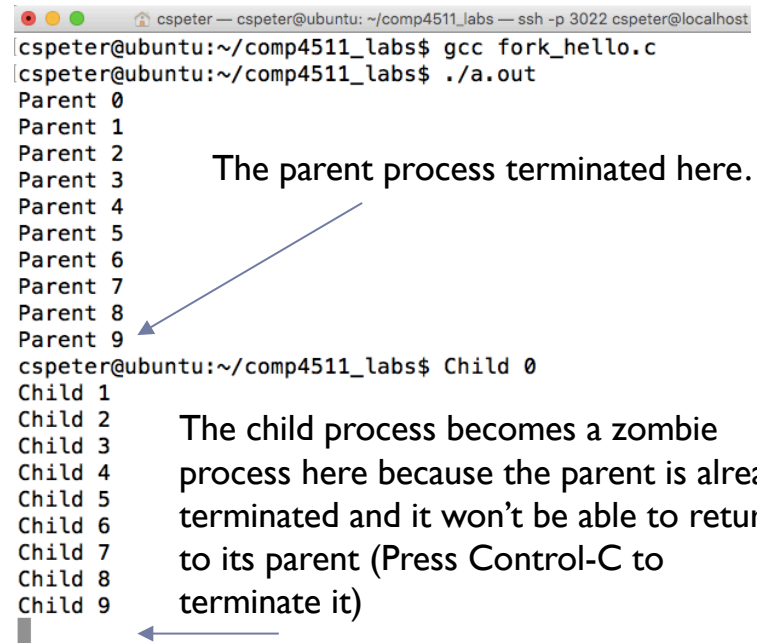
What is `fork()` good for?

- ▶ Two processes run concurrently if their flows overlap in time
- ▶ What does concurrency gain us?
 - ▶ The appearance that multi-actions are occurring at the same time
 - ▶ If done right, your program can improve throughput
- ▶ `fork()` creates a new process that runs concurrently
- ▶ Why concurrency?
 - ▶ Exploit natural concurrent structure of an application
 - ▶ Easier to program multiple independent and concurrent activities
 - ▶ Better resource utilization
 - ▶ Resource unused by one application can be used by others
 - ▶ Better average response time
 - ▶ No need to wait for other applications to make progress

fork HelloWorld - fork_hello.c

```
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>
int main() {
    pid_t pid;
    int i;
    /* Create a new process by
    duplicating the calling process */
    pid = fork();

    if ( pid > 0 ) {
        /* parent process */
        for (i=0; i<10; i++)
            printf("Parent %d\n",i);
    } else { /* child process */
        for (i=0; i<10; i++)
            printf("Child %d\n",i);
    }
    return 0;
}
```



```
cspeter — cspeter@ubuntu: ~/comp4511_labs — ssh -p 3022 cspeter@localhost
cspeter@ubuntu:~/comp4511_labs$ gcc fork_hello.c
cspeter@ubuntu:~/comp4511_labs$ ./a.out
Parent 0
Parent 1
Parent 2
Parent 3
Parent 4
Parent 5
Parent 6
Parent 7
Parent 8
Parent 9
cspeter@ubuntu:~/comp4511_labs$ Child 0
Child 1
Child 2
Child 3
Child 4
Child 5
Child 6
Child 7
Child 8
Child 9
```

The parent process terminated here...

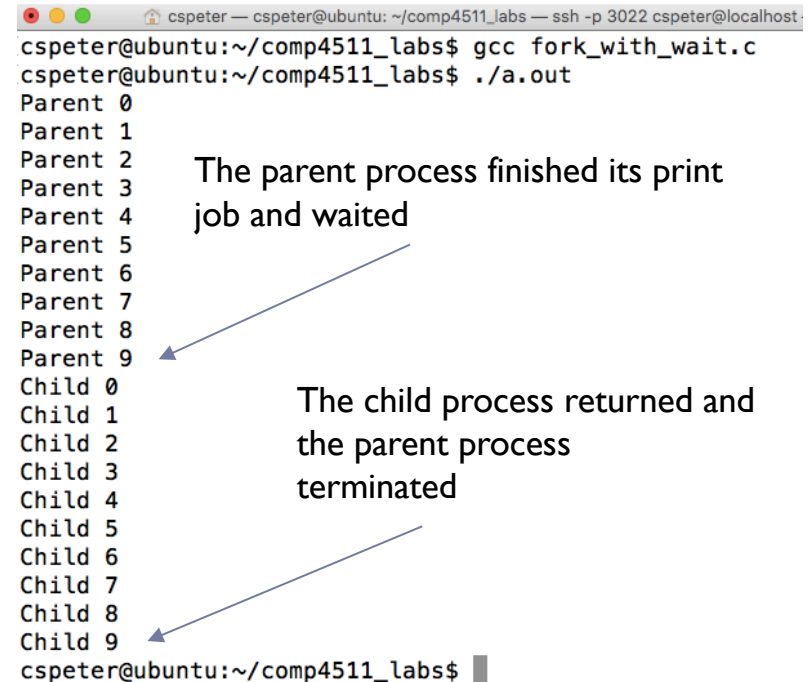
The child process becomes a zombie process here because the parent is already terminated and it won't be able to return to its parent (Press Control-C to terminate it)

Zombies

- ▶ What happens on termination?
 - ▶ When process terminates, still consumes system resources
 - ▶ Entries in various table & information maintained by OS
- ▶ Called a “zombie”, waiting parent to reap it
- ▶ What if parent does not reap?
 - ▶ If any parent terminates without reaping a child, then child will be reaped by `init` process

A better fork example – fork_with_wait.c

```
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>
int main() {
    pid_t pid;
    int i;
    /* Create a new process by
    duplicating the calling process */
    pid = fork();
    if ( pid > 0 ) { /* parent process */
        for (i=0; i<10; i++)
            printf("Parent %d\n",i);
        wait(0);
        /* Wait for the child process */
    } else { /* child process */
        for (i=0; i<10; i++)
            printf("Child %d\n",i);
    }
    return 0;
}
```



```
cspeter@ubuntu:~/comp4511_labs$ gcc fork_with_wait.c
cspeter@ubuntu:~/comp4511_labs$ ./a.out
Parent 0
Parent 1
Parent 2
Parent 3
Parent 4
Parent 5
Parent 6
Parent 7
Parent 8
Parent 9
Child 0
Child 1
Child 2
Child 3
Child 4
Child 5
Child 6
Child 7
Child 8
Child 9
cspeter@ubuntu:~/comp4511_labs$
```

The parent process finished its print job and waited

The child process returned and the parent process terminated

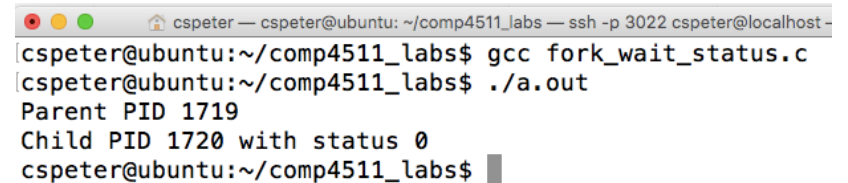
Waiting a child process: wait, waitpid, waitid

```
#include <sys/types.h>
pid_t wait(int *status);
pid_t waitpid(pid_t pid, int *status, int options);
int waitid(idtype_t idtype, id_t id, siginfo_t *infop, int
options);
```

- ▶ All of these system calls are used to wait for state changes in a child of the calling process
- ▶ status is the pointer pointing to an int that captures the return status. You can input 0 (NULL) if no return status is needed
- ▶ Reference: <http://linux.die.net/man/2/wait> OR man wait

Checking return status – fork_wait_status.c

```
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>
int main() {
    int child_status;
    pid_t child_pid ;
    pid_t pid = fork();
    if ( pid == 0 ) { /* child */
        return 0; }
    else { /* parent */
        printf("Parent PID %d\n", getpid());
        child_pid = wait(&child_status);
        printf("Child PID %d with status %d\n"
            , child_pid, child_status);
    }
    return 0;
}
```




```
cspeter — cspeter@ubuntu: ~/comp4511_labs — ssh -p 3022 cspeter@localhost -
cspeter@ubuntu:~/comp4511_labs$ gcc fork_wait_status.c
cspeter@ubuntu:~/comp4511_labs$ ./a.out
Parent PID 1719
Child PID 1720 with status 0
cspeter@ubuntu:~/comp4511_labs$
```

Status number 0 means exit without any error


Note: If parent has multiple children, wait will return when one of them (order not known!) completes.

Child process creation pattern: Chain and fan

Chain




Fan




Child process creation pattern: Chain and fan

Chain

```
pid_t childpid;  
for (i = 1; i < n; ++i)  
    if (childpid = fork())  
        break;
```




Fan



Child process creation pattern: Chain and fan


Chain

```
pid_t childpid;  
for (i = 1; i < n; ++i)  
    if (childpid = fork())  
        break;
```



Fan

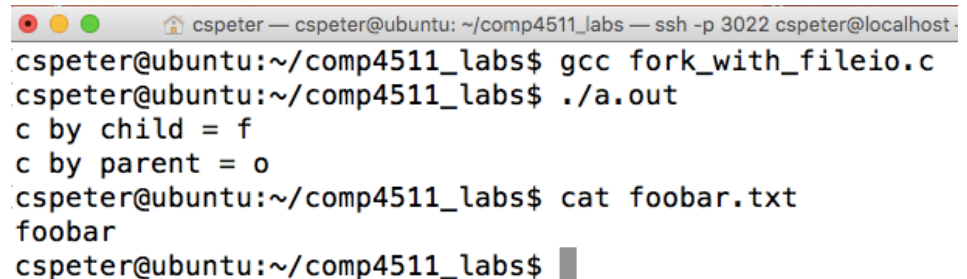
```
pid_t childpid;  
for (i = 1; i < n; ++i)  
    if ((childpid = fork()) <= 0)  
        break;
```



fork with fileI/O – child inherits open files

► The output of the program

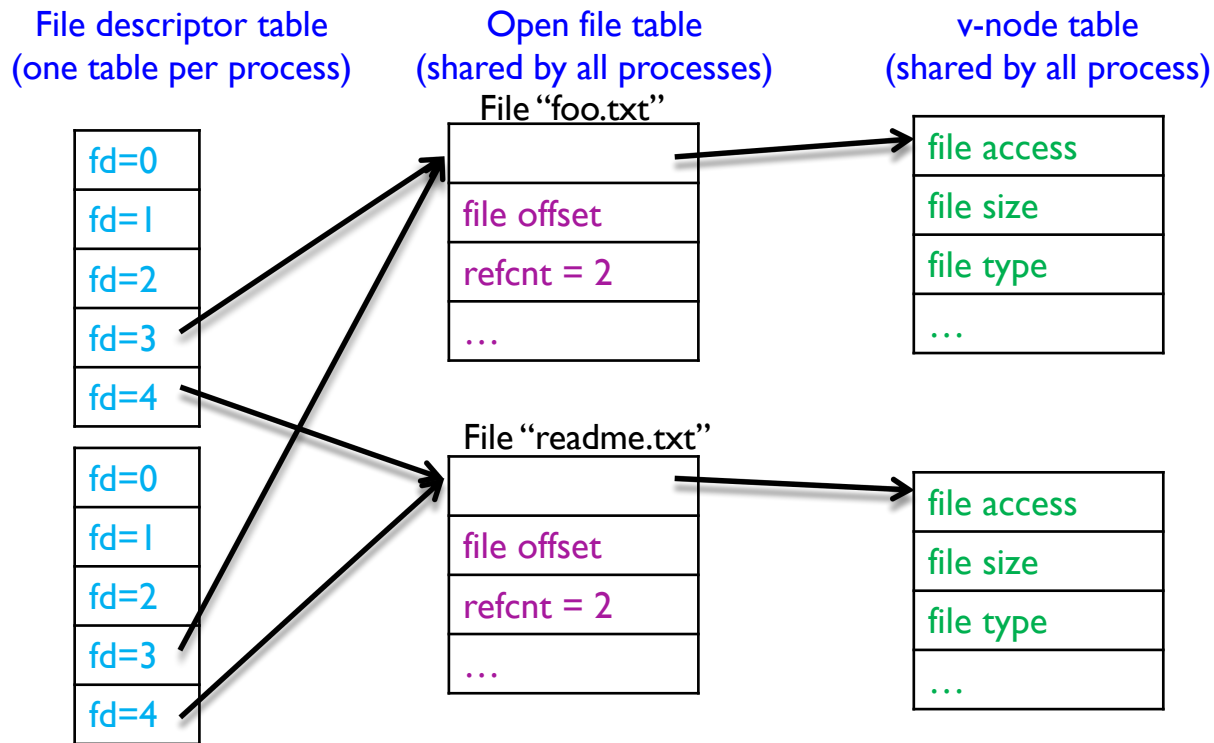
```
#include <stdio.h>
#include <unistd.h> /* POSIX header */
#include <fcntl.h> /* open for POSIX */
#include <sys/types.h>
/* foobar.txt: A text file contains 6 characters: foobar */
int main() {
    char c;
    int fd = open("foobar.txt", O_RDONLY, 0);
    pid_t pid = fork();
    if ( pid == 0 ) { /* child process */
        read(fd, &c, 1); /* read a char */
        printf("c by child = %c\n", c);
        return 0; /* terminate */
    }
    wait(0); /* wait for the child process */
    read(fd, &c, 1);
    printf("c by parent = %c\n", c);
    close(fd);
    return 0;
}
```



A terminal window titled 'cspeter — cspeter@ubuntu: ~/comp4511_labs — ssh -p 3022 cspeter@localhost'. The terminal shows the following commands and output:

```
cspeter@ubuntu:~/comp4511_labs$ gcc fork_with_fileio.c
cspeter@ubuntu:~/comp4511_labs$ ./a.out
c by child = f
c by parent = o
cspeter@ubuntu:~/comp4511_labs$ cat foobar.txt
foobar
cspeter@ubuntu:~/comp4511_labs$
```

Child process inherits open files



When a process terminates

- ▶ When a child process terminates:
 - ▶ Open files are flushed and closed
 - ▶ Child's resources are de-allocated
 - ▶ File descriptors, memory, semaphores, file locks, ...
- ▶ Parent process is notified via signal SIGCHLD
- ▶ Exit status is available to parent via `wait()`

Voluntary termination	Involuntary termination
Normal exit: <code>exit(0)</code>	Fatal error: divide by 0, core dump, segment fault
Error exit: <code>exit(1)</code>	Killed by another process: <code>kill()</code>

Sleep - sleep the current process for the specified number of seconds

```
#include <stdio.h>
#include <unistd.h>

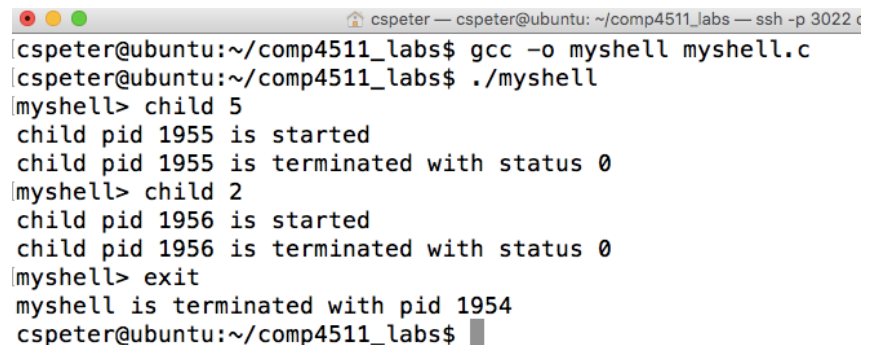
int main() {
    printf("Start to sleep for 10s\n");
    sleep(10);
    printf("End of sleep\n");
    return 0;
}
```

```
cspeter@ubuntu:~/comp4511_labs$ gcc sleep_example.c
cspeter@ubuntu:~/comp4511_labs$ ./a.out
Start to sleep for 10s
End of sleep
cspeter@ubuntu:~/comp4511_labs$
```

Lab Exercise: Create a simple shell with child process creation

- ▶ `myshell_skeleton.c` is provided as the starting point
- ▶ Complete the `process_cmd` function
- ▶ Make sure that you use the `wait` function to reap the child process
- ▶ Commands supported:
 - ▶ `exit`
 - ▶ `child [n]`
 - ▶ Create a child process that runs for `n` seconds (using `sleep`)
 - ▶ Print out the PID of the child process
 - ▶ Print out the status code when the child process terminates

```
void process_cmd(char *cmdline)
{
    // printf("%s\n", cmdline);
}
```



```
cspeter — cspeter@ubuntu: ~/comp4511_labs — ssh -p 3022 c
cspeter@ubuntu:~/comp4511_labs$ gcc -o myshell myshell.c
cspeter@ubuntu:~/comp4511_labs$ ./myshell
myshell> child 5
child pid 1955 is started
child pid 1955 is terminated with status 0
myshell> child 2
child pid 1956 is started
child pid 1956 is terminated with status 0
myshell> exit
myshell is terminated with pid 1954
cspeter@ubuntu:~/comp4511_labs$
```