

COMP 4511

Inter-Process Communication
Part 1: Pipe and FIFO

Inter-process Communication (IPC) in Linux

- ▶ There are a few methods which can communicate between among processes on the same host
 - ▶ Files
 - ▶ Pipes (will be discussed in this lab)
 - ▶ FIFOs (will be discussed in this lab)
 - ▶ Message Queues
 - ▶ Shared Memory (will be discussed in the next lab)
 - ▶ Unix domain sockets
- ▶ For synchronization on the same host
 - ▶ Signals (for synchronization)
 - ▶ Semaphores (for synchronization)
- ▶ For processes on separated hosts
 - ▶ Remote Procedure Call (RPC)
 - ▶ Socket

Pipes



Introduction

- ▶ Piping is a process where the output of one process is made as the input of another process
- ▶ Example: Piping in Shell (using the `|` operator)
 - ▶ `ps aux | grep root | sort | less`
 - ▶ The `ps aux` command will output a list of running processes and the corresponding information.
 - ▶ After that, the output will be treated as an input of `grep`, which is a program to match the given pattern `root`
 - ▶ The lines will be sorted in an alphabetical order using the `sort` command
 - ▶ The final text result will be displayed using the `less` command

Two types of pipes

▶ Unnamed pipe

- ▶ A unnamed pipe does not associated with any file
- ▶ Can only be shared by related processes (descendants of a process that creates the unnamed pipe)
- ▶ Pipe is an internal data structure maintained by the kernel
- ▶ Read and write processes are assumed to run concurrently

▶ Named pipe

- ▶ Like a file (create a named pipe, open, read/write)
- ▶ Can be shared by any process
- ▶ FIFO

Pipe in <stdio.h>

- ▶ popen, pclose

- ▶ **FILE** *popen(const char *command, const char *type);

- ▶ **int** pclose(FILE *stream);

- ▶ popen and pclose are higher level, invoking the shell to complete the operations

- ▶ We should **avoid using** popen and pclose for the programming assignments

- ▶ Example:

```
#include <stdio.h>
#define MAXSTRS 5
int main()
{
    int i;
    FILE *pipe_fp;
    char *strings[MAXSTRS] = {"Ann", "Dog", "Bob", "Egg", "Cat"};

    pipe_fp = popen("sort", "w"); // create a "sort" pipe
    if (!pipe_fp) {
        perror("popen"); // error handling
        exit(1);
    } else {
        for (i=0; i<MAXSTRS;i++) {
            fputs(strings[i], pipe_fp); // send a string to the pipe
            fputc('\n', pipe_fp); // send the endline character
        }
    }
    pclose(pipe_fp); // close the pipe
    return 0;
}
```

Unnamed pipe

```
#include <unistd.h>
int pipe(int pfd[2]);
```

- ▶ Create a message pipe
 - ▶ Anything can be written to the pipe, and read from the other end
 - ▶ Data is received in the order it was sent
 - ▶ OS enforces mutual exclusion: only one process at a time
 - ▶ Accessed by a file descriptor
 - ▶ Processes sharing the pipe must have the same parent
- ▶ Returns a pair of file descriptors (
 - ▶ `pfd[0]` is the read end
 - ▶ `pfd[1]` is the write end



Pipe example – Necessary header files

```
#include <stdio.h>
#include <stdlib.h>
#include <errno.h>
#include <sys/types.h>
#include <unistd.h>

int main()
{
    // The main program for pipes
}
```


Pipe example (Child => Parent)

```
int main()
{
    int pfd[2];
    char buf[30];
    pipe(pfd); /* Create a message pipe*/
    pid_t pid = fork(); /* 0 (child), non-zero (parent) */
    if ( pid == 0 ) {
        printf("CHILD: writing to pipe\n");
        close(pfd[0]);
        write(pfd[1], "test", 5);
        printf("CHILD: exiting\n");
    } else {
        printf("PARENT: reading from pipe\n");
        close(pfd[1]);
        read(pfd[0], buf, 5);
        wait(NULL); /* Wait until the child returns*/
        printf("PARENT: read \"%s\"\n", buf);
    }
    return 0;
}
```



Pipe example (Parent => Child)

```
int main()
{
    int pfd[2];
    char buf[30];
    pipe(pfd); /* Create a message pipe*/
    pid_t pid = fork(); /* 0 (child), non-zero (parent) */
    if ( pid != 0 ) {
        printf("PARENT: writing to pipe\n");
        close(pfd[0]);
        write(pfd[1], "test", 5);
        wait(NULL); /* Wait until the child returns*/
        printf("PARENT: exiting\n");
    } else {
        printf("CHILD: reading from pipe\n");
        close(pfd[1]);
        read(pfd[0], buf, 5);
        printf("CHILD: read \"%s\"\n", buf);
    }
    return 0;
}
```



Duplicating a file descriptor

```
#include <unistd.h>
int dup(int oldfd);
```

- ▶ Create a copy of an open file descriptor: put new copy in first **unused** file descriptor
- ▶ Returns
 - ▶ Return value ≥ 0 : success, returns new file descriptor
 - ▶ Return value = -1: error, check value of **errno**
- ▶ Parameters
 - ▶ **oldfd**: the open file descriptor to be duplicated

Example: A command-line pipe

ls | wc -l

Can we implement a command-line pipe
with `pipe()`?

How do we attach the stdout of `ls`
to the stdin of `wc`?


Command-line pipe: ls | wc -l

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
int main() {
    int pfd[2];
    pipe(pfd);
    pid_t pid = fork(); /* 0 (child), non-zero (parent) */
    if ( pid == 0 ) { /* The child process*/
        close(1); /* close stdout */
        dup(pfd[1]); /* make stdout as pipe input */
        close(pfd[0]); /* don't need this */
        execlp("ls", "ls", NULL);
    } else { /* The parent process*/
        close(0); /* close stdin */
        dup(pfd[0]); /* make stdin as pipe output */
        close(pfd[1]); /* don't need this */
        wait(0); /* wait for the child process */
        execlp("wc", "wc", "-l", NULL);
    }
    return 0;
}
```



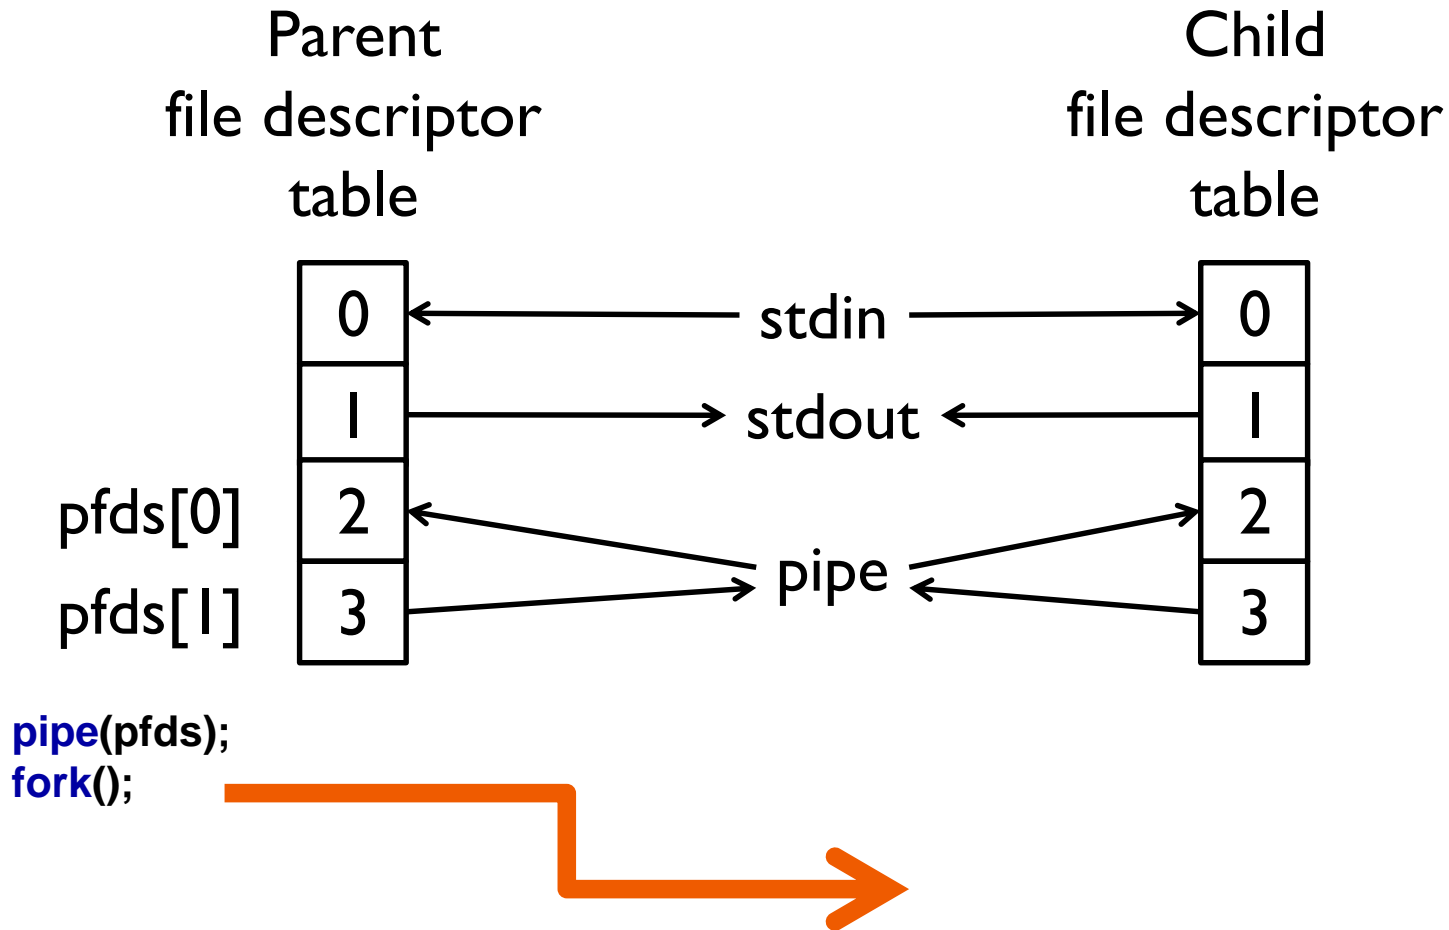
Command-line pipe: `ls | wc -l`

Parent
file descriptor
table

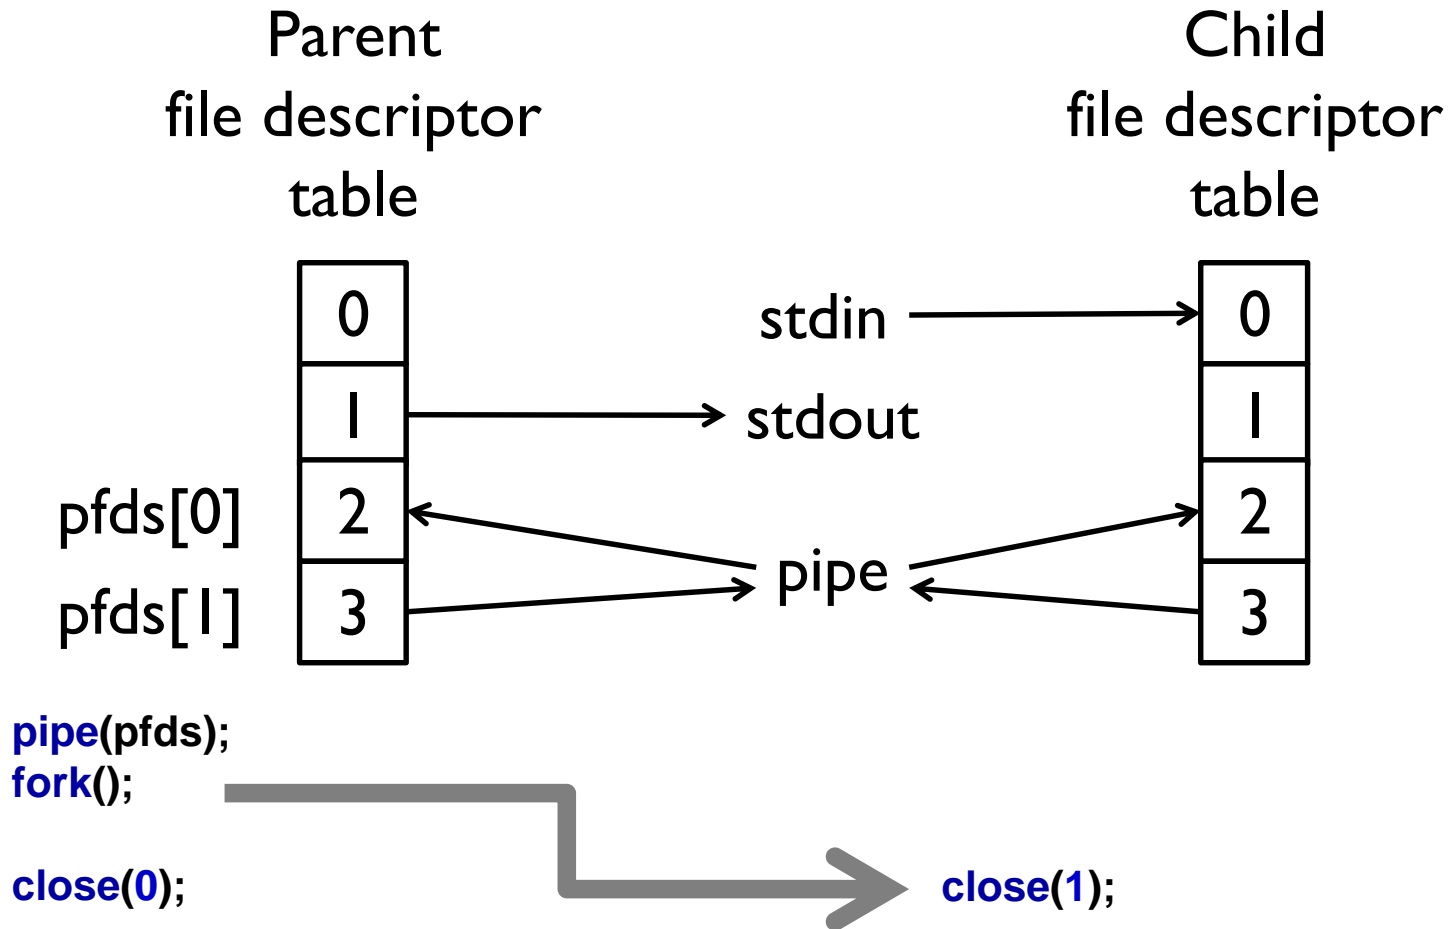


`pipe(pfd);`

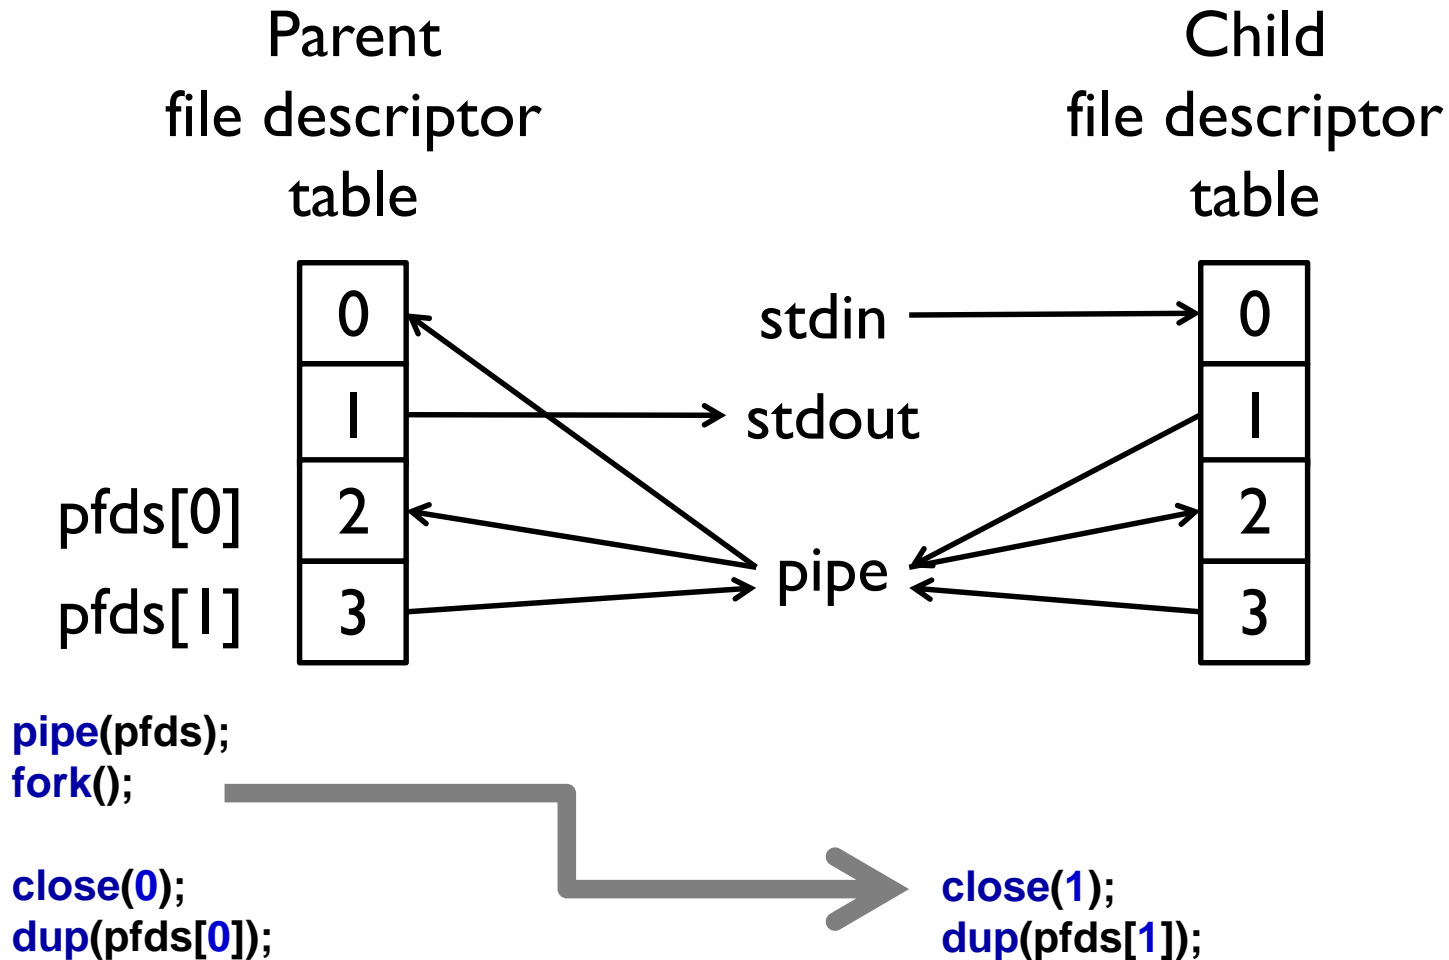
Command-line pipe: `ls | wc -l`



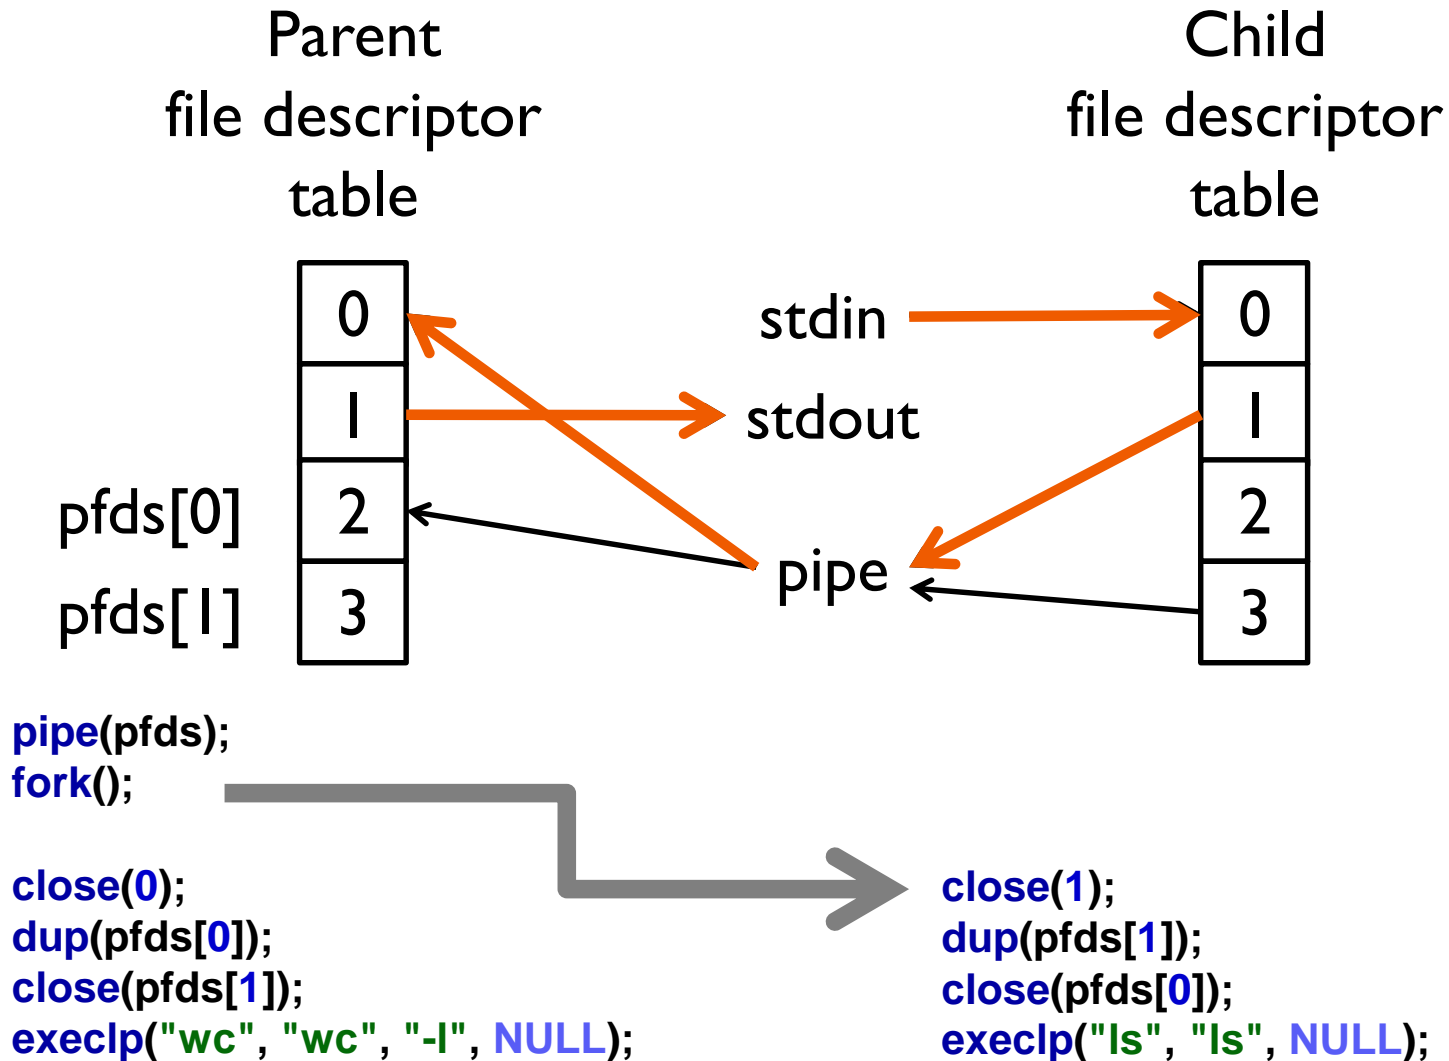
Command-line pipe: `ls | wc -l`



Command-line pipe: `ls | wc -l`



Command-line pipe: `ls | wc -l`



Duplicating a file descriptor

```
#include <unistd.h>
```

```
int dup2(int oldfd, int newfd);
```

- ▶ Create a copy of an open file descriptor: put new copy in specified location (...after closing **newfd**, if it was open)
- ▶ Returns
 - ▶ Return value ≥ 0 : success, returns new file descriptor
 - ▶ Return value $= -1$: error, check value of **errno**
- ▶ Parameters
 - ▶ **oldfd**: the open file descriptor to be duplicated

Command-line pipe: ls | wc -l (using dup2)

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
int main() {
    int pfd[2];
    pipe(pfd);
    pid_t pid = fork(); /* 0 (child), non-zero (parent) */
    if ( pid == 0 ) { /* The child process*/
        close(1); /* close stdout */
        dup2(pfd[1],1); /* make stdout as pipe input */
        close(pfd[0]); /* don't need this */
        execlp("ls", "ls", NULL);
    } else { /* The parent process*/
        close(0); /* close stdin */
        dup2(pfd[0],0); /* make stdin as pipe output */
        close(pfd[1]); /* don't need this */
        wait(0); /* wait for the child process */
        execlp("wc", "wc", "-l", NULL);
    }
    return 0;
}
```

Lab exercise: Add pipe support to the simple shell

- ▶ Based on the simple shell exercise (`myshell_skeleton.c`) in the previous lab, implement a pipe operation:

- ▶ `myshell> [command A] | [command B]`

- ▶ In the current stage, you are **NOT** required to support multi-level pipe as follows:

- ▶ `myshell> [command A] | [command B] | [command C]`

FIFOs

FIFOs

- ▶ A pipe disappears when no process has it open
- ▶ FIFOs = **named pipes**
 - ▶ Special pipes that persist even after all processes have closed them
 - ▶ Actually implemented as a file

```
#include <sys/types.h>
```

```
#include <sys/stat.h>
```

```
/* Somewhere inside your main program.. */
```

```
int status;
```


```
status = mkfifo("/home/cnd/mod_done",
```

```
    /* mode=0644 as the second parameter */
```

```
    S_IWUSR | S_IRUSR | S_IRGRP | S_IROTH);
```

Communication over of FIFO

- ▶ First **open** blocks until second process opens the FIFO
- ▶ Can use **O_NONBLOCK** flag to make operations non-blocking
- ▶ FIFO is persistent: can be used multiple times
- ▶ Like pipes, OS ensures atomicity of writes and reads



FIFOs

- ▶ It can be used between unrelated processes
- ▶ When a FIFO is created, it must be opened with syscall `open()` or stream operations such as `fopen()`, but you cannot use `lseek()`
- ▶ You cannot open FIFO with “read+write” mode
- ▶ It has to be open at both ends simultaneously before you can proceed to do any input or output operations on it

```
/* server */
writefd = open(FIFO1, O_WRONLY);
readfd = open(FIFO2, O_RDONLY);

/* client */
readfd = open(FIFO1, O_RDONLY);
writefd = open(FIFO2, O_WRONLY);
```

Can we change the order of open sequences?

FIFO example: producer-consumer

- ▶ **Producer:**
 - ▶ Writes to FIFO
- ▶ **Consumer:**
 - ▶ Reads from FIFO
 - ▶ Outputs data to file
- ▶ FIFO ensures atomicity of write

FIFO example

```
#include <errno.h>
#include <fcntl.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/stat.h>
int main (int argc, char *argv[])
{
    int requestfd;

    if (argc != 2) {
        /* name of consumer fifo on the command line */
        fprintf(stderr, "Usage: %s fifoname > logfile\n", argv[0]);
        return 1;
    }

    /* Note: Continue in the next slide! */
```



FIFO example

```
/* create a named pipe to handle incoming requests */
if ((mkfifo(argv[1], S_IRWXU | S_IWGRP | S_IWOTH) == -1)
    && (errno != EEXIST)) {
    perror("Server failed to create a FIFO");
    return 1;
}

/* open a read/write communication endpoint to the pipe */
if ((requestfd = open(argv[1], O_RDWR)) == -1) {
    perror("Server failed to open its FIFO");
    return 1;
}

/* The remaining part of the program
   Write to pipe like you would to a file */
```